

Discrete Flow Matching: Tutorial and Code

May 28, 2025

Introduction

This tutorial bridges the gap between the rigorous mathematical formulation of *Discrete Flow Matching* for continuous time Markov chains and a minimal, practical PyTorch implementation. We begin by defining a factorized mixture path between a simple source distribution and an observed target distribution, derive the coordinatewise CTMC velocities via the Kolmogorov forward equation, and show how to assemble those into a learnable model. Each step of the derivation is followed by the code that implements it. You will find:

- A concise statement of each mathematical object (probability path, Bregman divergence, CTMC velocity, posterior parameterization).
- The exact lines of code hidden helpers showing how to (1) sample the mixture path, (2) compute the matching loss, and (3) run the coordinatewise Euler sampler.
- Tips on how to factorize over dimensions so that both training and sampling remain tractable even in high dimensions.

The basic idea is straightforward:

1. Sample a data point

$$X_1 \sim q_{\text{data}}.$$

This is your target (X_i^1) in every coordinate.

2. Sample a source

$$X_0$$

from your simple prior (e.g. uniform or noise).

3. Pick a time

$$t \sim U[0, 1].$$

4. Draw the intermediate state

$$X_t$$

from the mixture path

$$p_{t|0,1}(x \mid X_0, X_1) = \prod_{i=1}^D \left[\kappa(t) \delta(x_i, X_i^1) + (1 - \kappa(t)) \delta(x_i, X_i^0) \right].$$

In other words, each coordinate i of X_t is X_i^1 with probability $\kappa(t)$, or X_i^0 otherwise.

5. Compute the loss using the known true generator (which depends on X_i^1) and your models posterior

$$p_{1|t}^{\theta,i}(\cdot \mid X_t).$$

Because X_i^1 is always the data, it drives the true velocity in every coordinate, and X_t is exactly where you evaluate both true and model velocities.

$p_{1|t}^{\theta,i}(\cdot \mid X_t)$ is learned through a simple crossentropy objective:

$$L(\theta) = -\mathbb{E}_{t, X_0, X_1, X_t} \sum_{i=1}^d \log p_{1|t}^{\theta,i}(X_1^i \mid X_t) + \text{const.}$$

During inference, you reverse this:

1. Start from X_0 drawn from the prior.
2. Iteratively step forward in t using the learned posterior $p_{1|t}^{\theta,i}$ to sample each coordinates jump, until you reach X_1 your generated sample.

Note: X_i^1 in the formulas is always the data's coordinate i . X_t during training is sampled via the chosen mixture path $p_{t|0,1}$, ensuring the model learns to match that path's velocity field.

1 Factorized SourceSink Velocities

Because both the probability path and its velocity factorize over coordinates, we may treat each dimension independently. Writing a full state as $x = (x_1, \dots, x_d)$ and a proposed next state as $y = (y_1, \dots, y_d)$, we have

$$q_t(x) = \prod_{i=1}^d q_t^i(x_i),$$

and the full rate vector

$$u_t(y, x) = \sum_{i=1}^d \underbrace{\delta(y_i, x_i)}_{\prod_{j \neq i} \delta(y_j, x_j)} u_t^i(y_i, x).$$

Here

$$\delta(y_i, x_i) = \prod_{j \neq i} \delta(y_j, x_j)$$

ensures that only coordinate i is allowed to move, and

$$u_t^i(y_i, x) \in \mathbb{R}$$

is exactly the *sourcesink rate* for changing coordinate i from its current value x_i (the source) to a new value y_i (the sink). By factorizing both q_t and u_t this way, we reduce the high-dimensional CTMC rate problem into d independent scalar rate functions u_t^i .

2 Mixture-path Construction

Introduce an auxiliary coupling $Z = (X_0, X_1)$. For each coordinate $i \in \{1, \dots, d\}$, define the conditional marginal path

$$p_{t|0,1}^i(x_i | x_0, x_1) = \kappa(t) \delta(x_i, x_{1,i}) + (1 - \kappa(t)) \delta(x_i, x_{0,i}),$$

where we choose the simple scheduler

$$\kappa(t) = t, \quad \dot{\kappa}(t) = 1.$$

Under this choice, at time t each coordinate X_t^i equals its sink value X_1^i with probability t , and its source value X_0^i with probability $1 - t$. In other words,

$$\Pr[X_t^i = X_1^i | X_0 = x_0, X_1 = x_1] = t, \quad \Pr[X_t^i = X_0^i | X_0 = x_0, X_1 = x_1] = 1 - t.$$

Thus the full factorized path is

$$p_{t|0,1}(x | x_0, x_1) = \prod_{i=1}^d p_{t|0,1}^i(x_i | x_0, x_1).$$

3 Posterior-driven SourceSink Velocities

Once we have the pure sourcesink rates

$$u_t^i(y_i, x | x_0, x_1) = \frac{1}{1-t} \left[\delta(y_i, x_{1,i}) - \delta(y_i, x_i) \right]$$

for the mixture path, we introduce a learnable percoordinate targetinitial posterior

$$p_{\theta, i, 1|t}(x_{1,i} | x) \approx p_{i, 1|t}(x_{1,i} | x).$$

We then assemble the marginal velocity as

$$u_t^i(y_i, x) = \sum_{x_{1,i}} u_t^i(y_i, x | x_0, x_1) p_{\theta, i, 1|t}(x_{1,i} | x) = \frac{1}{1-t} \sum_{x_{1,i}} \left[\delta(y_i, x_{1,i}) - \delta(y_i, x_i) \right] p_{\theta, i, 1|t}(x_{1,i} | x).$$

Plugging this into the Eulerstep transition kernel

$$P(X_{t+h,i} = y_i \mid X_t = x) = \delta(y_i, x_i) + h u_t^i(y_i, x) + o(h)$$

yields

$$P(X_{t+h,i} = y_i \mid X_t = x) = \sum_{x_{1,i}} \left[\delta(y_i, x_i) + \frac{h}{1-t} (\delta(y_i, x_{1,i}) - \delta(y_i, x_i)) \right] p_{\theta,i,1|t}(x_{1,i} \mid x) + o(h).$$

In practice, sampling is implemented coordinatewise by:

1. Draw $x_{1,i} \sim p_{\theta,i,1|t}(\cdot \mid x)$.
2. With probability $\frac{h}{1-t}$, set $X_{t+h,i} = x_{1,i}$; otherwise keep $X_{t+h,i} = x_i$.

This makes explicit how the learned posterior per coordinate drives the sourcesink transitions during sampling.

4 Coordinatewise Euler Sampling

For a small time step h , the CTMC transition kernel factorizes over coordinates:

$$P(X_{t+h,i} = y_i \mid X_t = x) = \delta(y_i, x_i) + h u_t^i(y_i, x) + o(h).$$

Substituting our posteriorassembled velocity

$$u_t^i(y_i, x) = \frac{1}{1-t} \sum_{x_{1,i}} [\delta(y_i, x_{1,i}) - \delta(y_i, x_i)] p_{\theta,i,1|t}(x_{1,i} \mid x),$$

we obtain

$$P(X_{t+h,i} = y_i \mid X_t = x) = \sum_{x_{1,i}} \left[\delta(y_i, x_i) + \frac{h}{1-t} (\delta(y_i, x_{1,i}) - \delta(y_i, x_i)) \right] p_{\theta,i,1|t}(x_{1,i} \mid x) + o(h).$$

Hence the coordinatewise sampling algorithm is:

1. Draw $x_{1,i} \sim p_{\theta,i,1|t}(\cdot \mid x)$.
2. With probability $\frac{h}{1-t}$, set $X_{t+h,i} = x_{1,i}$; otherwise keep $X_{t+h,i} = x_i$.

5 PerCoordinate Posterior Parameterization and Training Loss

We parameterize the marginal targetinitial posterior per coordinate

$$p_{\theta,i,1|t}(x_{1,i} \mid x_t)$$

via a neural network (embedded and flattened as in Section ??). By Proposition 1 the Conditional Discrete Flow Matching loss over our mixture path reduces to a sum of generalizedKL Bregman divergences per coordinate:

$$L(\theta) = \mathbb{E}_{t, X_0, X_1, X_t} \left[\sum_{i=1}^d D_{\text{gKL}} \left(u_t^i(\cdot, x_t \mid x_0, x_1), u_t^{\theta, i}(\cdot, x_t) \right) \right].$$

For the mixturepath generator

$$u_t^i(y_i, x_t \mid x_0, x_1) = \frac{1}{1-t} [\delta(y_i, x_{1,i}) - \delta(y_i, x_{t,i})],$$

choosing D_{gKL} yields the simple crossentropy objective

$$L(\theta) = -\mathbb{E}_{t, X_0, X_1, X_t} \sum_{i=1}^d \log p_{\theta, i, 1|t}(X_{1,i} \mid X_t) + \text{const.}$$

In code, using the librarys generalizedKL loss wrapper:

```
from flow_matching.loss import MixturePathGeneralizedKL

loss_fn = MixturePathGeneralizedKL(path=path)

for x_0, x_1 in dataloader:
    t = torch.rand(batch_size) * (1.0 - 1e-3)
    sample = path.sample(t=t, x_0=x_0, x_1=x_1)
    logits = model(sample.x_t, sample.t)
    loss = loss_fn(
        logits=logits,
        x_1=sample.x_1,
        x_t=sample.x_t,
        t=sample.t
    )
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```